

# CSE 5441 (Fall 2019, Dr. Jones)

## CUDA (Lab 4)

Caleb Lehman  
[lehman.346@osu.edu](mailto:lehman.346@osu.edu)

November 16, 2019

## Overview

For this lab, I parallelized my serial C program to perform Adaptive Mesh Refinement (AMR)<sup>1</sup> using the CUDA API. I created 2 programs, `cuda-disposable` and `cuda-persistent`, which mirror the disposable and persistent threads models used in the previous labs. In particular, the `disposable` code runs a new kernel for each iteration, while the `persistent` code has a single kernel which synchronizes the threads between iterations<sup>2</sup>.

## Tests

### Environment

The program was developed and tested on the [Owens cluster](#) at the [Ohio Supercomputer Center](#). Note that all my other labs were developed and tested on the [Pitzer cluster](#), which has different specifications.

For development and testing, I loaded the `cuda/9.2.88` module, which allowed the program to be compiled with version 9.2.88 of the `nvcc` compiler, as well as setting some environment variables pointing to CUDA-related headers and libraries.

For testing, I loaded the `python/3.6-conda5.2` module, which loads a python environment with the NumPy, SciPy, and Matplotlib packages, among others. Python is only necessary for collecting and plotting the data from testing, not for the actual execution of the program.

### Timing

I collected timing data using the same 4 methods as the first several labs: `time`, `clock`, and `clock_gettime` from the `"time.h"` header, and the UNIX utility `time`.

As with the second and third labs, I didn't use the results from the `clock` function from the `"time.h"` header, since it reports CPU time, not wall time. The other methods all returned values within 1 second of each other. The `time` function declared in `"time.h"` returns an integer number of seconds, but the other two methods return with sub-second precision. *For consistency, I used the `clock_gettime` function for all results in this report.*

### Test Files

Dr. Jones provided the `testgrid_400.12206` test file. As part of lab 1, I reduced the  $\alpha$  (affect rate) and  $\varepsilon$  parameters until the serial runtime increased into the 3 to 6 minute range. In particular, I selected  $\alpha = 0.01$  and  $\varepsilon = 0.02$ , for which the serial program completed in 261 seconds. However, my CUDA implementations yielded very poor performance, so I ended up using parameters of  $\alpha = 0.1$  and  $\varepsilon = 0.1$ <sup>3</sup>. For this reason, and due to using the Owens cluster instead of the Pitzer cluster, it isn't particularly useful to compare the results from this lab to the results from labs 2 and 3.

<sup>1</sup>See lab 1 or project descriptions for details about AMR computation.

<sup>2</sup>Because the CUDA API doesn't directly provide a simple way to synchronize blocks, the `cuda-persistent` version is limited to running with a single block of threads.

<sup>3</sup>In any case, these parameters really only affect the final DSV values and number of iterations, not the actual amount of work done per iteration, so most of the relative results found in this report should be the same for different values of  $\alpha$ ,  $\varepsilon$ .

## Results

The output was consistent across both programs and was as follows:

- `testgrid_400_12206`: 75197 iterations,  $(max, min) = (0.077873, 0.086525)$

The relevant runtimes for the serial, `cuda-persistent`, and `cuda-disposable` versions on the test case are visualized in the figures below. I found that the `cuda-persistent` version had optimal performance around 320 threads (and 1 block) and the `cuda-disposable` version had optimal performance with 32 threads and 10 blocks. Those values are now hardcoded into the program, as requested. With the optimal values for threads and blocks, `cuda-persistent` ran in 106 seconds and `cuda-disposable` ran in 60 seconds.

The two biggest implications I see in this experiment are:

- My serial program (with the same  $\alpha = 0.1, \varepsilon = 0.1$  parameters) runs in 14 seconds, so the CUDA versions are significantly slower. Since other parallelization methods (`OpenMP` and `pthread`s) improved performance, this indicates that my CUDA implementation was probably poorly done. I wasn't able to determine why the performance was so bad, but I suspect that it has something to do with cache performance, since the cache on the GPU is much smaller than that on the CPU.
- When comparing the `cuda-persistent` and `cuda-disposable` version with the same thread/block parameters, the `cuda-disposable` version consistently performed better. For example, when both programs used 256 threads and 1 block, `cuda-disposable` ran in 110 seconds, while `cuda-persistent` ran in 124 seconds. This shows that it is faster to kill and restart the entire kernel (`cuda-disposable`) than to just synchronize the threads (which do slightly uneven amounts of work in `cuda-persistent`). This is the opposite of the behaviour I noticed in the `pthread`s lab, confirming that CUDA threads are much more lightweight than even the relatively lightweight `pthread`s.




Figure 1: The runtimes of the `cuda-persistent` (left) and `cuda-disposable` (right) parallelizations. Note that, due to the poor performance of the CUDA parallelization, higher  $\alpha$  and  $\varepsilon$  parameters were used compared to previous labs. For comparison, my serial program runs in 14 seconds with these parameters. **Important note:** the `cuda-persistent` plot displays the runtimes on a scale starting at 100 seconds.

The following is the series of computations performed for the  $i$ th box during each iteration, which make up basically all of the “useful” work done in the programs:

```
updated_vals[i] = box->self_overlap * current_vals[i];
for (int nhbr = 0; nhbr < box->num_nhbrs; ++nhbr) {
    updated_vals[i] += box->overlaps[nhbr] * current_vals[box->nhbr_ids[nhbr]];
}
updated_vals[i] /= box->perimeter;
updated_vals[i] = current_vals[i] * (1 - affect_rate)
    + updated_vals[i] * affect_rate;
```

The number of arithmetic operations in the above sample is  $1 + 2 \cdot N_i + 1 + 4$ , where  $N_i$  is the number of neighbors of the  $i$ th box. Letting  $n$  be the number of boxes,  $N$  be the total number of neighbors (counted with multiplicities), and  $I$  be the number of iterations, the total number of arithmetic operations is given by  $I \cdot (6n + 2N)$ . For the given test file (`testgrid.400_12206`) and choice of parameters ( $\alpha = 0.1, \varepsilon = 0.1$ ), we have the values  $I = 75197$ ,  $n = 12206$ , and  $N = 71890$ . Substituting in these values, I computed  $GFlops = \frac{ops/10^9}{\#ofseconds} = \frac{16.32}{\#ofseconds}$  for each of the programs to get the following results<sup>4</sup>:

	serial	cuda-disposable	cuda-persistent
GFlops	1.166	0.272	0.154

## Project Usage

### Building

To build the `cude-disposable` and `cude-persistent` executables, navigate to the top level of the submitted directory and build as follows:

```
# Ensure that you have nvcc compiler

# Note that the provided makefile
# assumes the CUDA_HOME environment
# variable is set appropriately

# On the OSC clusters, this can be achieved
# with the command:
$ module load cuda

$ make
$ ls
... cuda-persistent cuda-disposable ...
```

### Running

The syntax to run the program is:

```
$ ./[program] [affect-rate] [epsilon] <[test-file]
```

where `program` is one of `{cuda-persistent, cuda-disposable}`

---

<sup>4</sup>I used the optimal times from each of the **CUDA** programs for this computation.